

AI-Powered Carbon Footprint Monitoring and Intelligent Decision-Support Platform

Complete Implementation Plan for 2-Semester Major Project

Executive Summary

Project Name: CarbonSense AI - Multi-Agentic Carbon Intelligence Platform

Vision: Build a production-grade, enterprise-ready platform that provides real-time carbon footprint monitoring, predictive analytics, and intelligent decision support using cutting-edge AI/ML techniques across multiple domains.

Core Philosophy: Not a prototype, not a dashboard - a scalable, secure, production-ready platform that enterprises and individuals can actually deploy.

1. PROBLEM ANALYSIS & RESEARCH GAP IDENTIFICATION

1.1 Critical Gaps Identified from Literature Survey

After analyzing 16 research papers, we identified the following critical gaps:

1. **Fragmented Solutions:** Existing systems focus on specific sectors (power plants, agriculture, factories) but lack unified, cross-domain monitoring
2. **Limited Real-World Deployment:** Most papers propose methodologies without real production implementation
3. **Weak Explainability:** AI models lack transparency, making them unsuitable for compliance and auditing
4. **No Personalization:** One-size-fits-all approach without behavioral adaptation
5. **Security & Trust Deficit:** Minimal focus on data privacy, authentication, and blockchain-based verification
6. **Reactive vs Proactive:** Systems monitor but don't intelligently predict and prevent
7. **Data Quality Issues:** Heavy reliance on manual data entry or incomplete datasets
8. **No Multi-Modal Integration:** Lack of computer vision, NLP, and time-series integration
9. **Missing Behavioral Psychology:** No integration of nudge theory and behavioral economics

10. **Lack of Continuous Learning:** Static models without online learning capabilities
-

2. PROPOSED SOLUTION: CarbonSense AI

2.1 Platform Overview

CarbonSense AI is a multi-agentic, AI-powered carbon intelligence platform that provides:

- **Real-time monitoring** across personal, organizational, and industrial carbon footprints
 - **Multi-modal AI agents** specializing in different emission sources
 - **Predictive analytics** using ensemble deep learning
 - **Explainable recommendations** with XAI techniques
 - **Behavioral nudging** using reinforcement learning
 - **Blockchain-verified** carbon credit tracking
 - **Computer vision-based** invoice/receipt processing
 - **NLP-powered** document analysis and policy extraction
 - **Secure, auditable** architecture with zero-trust security
-

3. UNIQUE VALUE PROPOSITIONS (USPs)

3.1 Core USPs That Differentiate Our Platform

1. **Multi-Agentic Architecture**
 - Specialized AI agents for different emission domains (transport, energy, food, waste, procurement)
 - Agents communicate and coordinate using a central orchestrator
 - Each agent uses domain-specific models and reasoning
2. **Computer Vision Integration**
 - Receipt/invoice OCR for automatic expense categorization
 - Image-based carbon estimation from uploaded bills
 - Visual product carbon labeling recognition
 - Satellite imagery analysis for organizational footprint (using public APIs)
3. **Advanced NLP Capabilities**
 - Extract carbon data from financial documents, reports, and policies
 - Sentiment analysis of sustainability reports
 - Automatic policy compliance checking
 - Conversational AI for user queries
4. **Explainable AI (XAI) First**
 - SHAP, LIME, and Attention mechanisms for all predictions
 - Visual explanations for every recommendation

- Trust scores for predictions
 - Compliance-ready audit trails
 - 5. **Behavioral AI & Gamification**
 - Reinforcement learning-based personalized nudges
 - Psychological profiling for intervention strategies
 - Social comparison with privacy preservation
 - Achievement system with carbon credits
 - 6. **Blockchain-Verified Carbon Credits**
 - Smart contracts for carbon offset verification
 - Immutable audit trails
 - Tokenized carbon credits
 - Decentralized verification network
 - 7. **Hybrid Intelligence**
 - Combines multiple deep learning architectures
 - Ensemble methods for robust predictions
 - Online learning for continuous improvement
 - Transfer learning across domains
 - 8. **Enterprise-Grade Security**
 - Zero-trust architecture
 - End-to-end encryption
 - Role-based access control (RBAC)
 - GDPR/privacy compliance
 - Secure multi-party computation for data sharing
 - 9. **Real-Time Processing**
 - Streaming data pipelines
 - Edge computing for IoT integration
 - Sub-second inference times
 - Scalable microservices architecture
 - 10. **Comprehensive Multi-Modal Integration**
 - Time-series (energy consumption patterns)
 - Text (documents, reports, receipts)
 - Images (bills, product labels)
 - Structured data (databases, APIs)
 - Semi-structured (JSON, XML from IoT)
-

4. SYSTEM ARCHITECTURE

4.1 High-Level Architecture

PRESENTATION LAYER

Web App (React)	Mobile App (React N)	Admin Dashboard	API Gateway
--------------------	-------------------------	--------------------	----------------

API GATEWAY LAYER
(Kong/AWS API Gateway + Rate Limiting)

AUTHENTICATION & SECURITY

Auth0/ Keycloak	JWT + OAuth2.0	Encryption Service
--------------------	-------------------	-----------------------

MICROSERVICES LAYER

User Service	Carbon Calculation	Agent Orchestrator	Analytics Service
NLP Service	Computer Vision	Time Series	Blockchain Service
Notification Service	Reporting Service	XAI Service	

AI/ML AGENT LAYER

Transport Agent	Energy Agent	Food Agent	Waste Agent
Procurement Agent	Behavioral Agent	Prediction Agent	XAI Agent

DATA LAYER

PostgreSQL (Relational)	MongoDB (Documents)	Redis (Cache)	TimescaleDB (TimeSeries)
S3/ MinIO (File Store)	Vector DB (Pinecone/ Qdrant)	Blockchain Ledger (Hyperledger)	

INFRASTRUCTURE LAYER

Kubernetes (Orchestr.)	Docker (Containers)	Apache Kafka	Airflow (Pipeline)
Prometheus/ Grafana	ELK Stack	MLflow (ML Ops)	

4.2 Multi-Agent AI Architecture

AGENT ORCHESTRATOR (Central Coordination & Routing)

- Agent Discovery & Registration
- Task Distribution & Load Balancing
- Inter-Agent Communication (Message Passing)
- Conflict Resolution & Decision Aggregation

↓	↓	↓	↓
TRANSPORT AGENT	ENERGY AGENT	FOOD AGENT	WASTE AGENT
• Route Optimization • Modal Split Analysis • EV vs ICE Comparison	• Appliance Detection • Consumption Prediction • Peak/Off-Peak Shift • Renewable	• Meal Recognition • Ingredient Analysis • Diet Optimization • Local vs	• Waste Classification • Recycling Potential • Composting Detection • Circular

Models: - LSTM - Transformer - RL (PPO)	Integration Models: - CNN-LSTM - GRU - Attention	Imported Models: - ResNet-50 - EfficientNet - BERT	Economy Models: - YOLO v8 - Random Forest
PROCUREMENT AGENT	BEHAVIORAL AGENT	PREDICTION AGENT	XAI AGENT
• Invoice OCR • Product Carbon DB • Supplier Ranking • Alternative Suggestions	• User Profiling • Nudge Generation • Gamification • Social Comparison • Achievement Tracking	• Time Series Forecasting • Anomaly Detection • Scenario Analysis • Risk Assessment	• SHAP • LIME • Attention Visualization • Feature Importance • Decision Trees • Counterfact. Explanation
Models: - Tesseract - BERT - EasyOCR - TF-IDF - Sentence Transformers	Models: - DQN (RL) - SARSA - Bandit Algorithms	Models: - Prophet - N-BEATS - Temporal Fusion Transform.	Framework: - SHAP - Captum - InterpretML

5. COMPREHENSIVE TECH STACK

5.1 Frontend Technologies

Component	Technology	Justification
Web Framework	React 18.x with TypeScript	Type safety, component reusability, excellent ecosystem
State Management	Redux Toolkit + RTK Query	Centralized state, automatic caching, optimistic updates
UI Framework	Material-UI (MUI) v5 + Tailwind CSS	Professional components + utility-first styling
Data Visualization	D3.js + Recharts + Plotly	Interactive charts, 3D visualizations, custom graphs

Component	Technology	Justification
3D Visualization	Three.js	Carbon footprint 3D visualizations
Maps	Mapbox GL JS / Leaflet	Route visualization, location-based emissions
Forms	React Hook Form + Yup	Efficient validation, better UX
Testing	Jest + React Testing Library	Unit and integration testing
Build Tool	Vite	Fast development, optimized builds
Mobile	React Native (optional for Sem 2)	Code reuse from web app

5.2 Backend Technologies

Component	Technology	Justification
Runtime	Node.js 20.x LTS	Non-blocking I/O, excellent for microservices
Framework	Express.js + Fastify	RESTful APIs + high-performance endpoints
API Gateway	Kong / AWS API Gateway	Rate limiting, authentication, routing
Authentication	Auth0 / Keycloak + JWT	OAuth2.0, SSO, multi-factor authentication
GraphQL	Apollo Server	Efficient data fetching for complex queries
WebSockets	Socket.io	Real-time notifications and updates
API Documentation	Swagger/OpenAPI 3.0	Auto-generated docs, testing interface
Validation	Joi / Zod	Schema validation
Testing	Jest + Supertest	API testing

5.3 AI/ML Technologies

Component	Technology	Justification
Primary Framework	PyTorch 2.x	Flexibility, research-friendly, dynamic graphs
Secondary Framework	TensorFlow 2.x / Keras	Production deployment, TF Serving

Component	Technology	Justification
Computer Vision	OpenCV + PIL + Torchvision	Image processing, model utilities
OCR	Tesseract + EasyOCR + PaddleOCR	Multi-language invoice/receipt processing
NLP	Transformers (Hugging Face)	BERT, GPT variants, sentence transformers
Time Series	Prophet + statsmodels + sktime	Forecasting, decomposition
XAI	SHAP + LIME + Captum + InterpretML	Model interpretability
RL	Stable-baselines3 + Ray RLlib	Reinforcement learning for behavioral agent
AutoML	Optuna + Ray Tune	Hyperparameter optimization
ML Ops	MLflow + DVC	Experiment tracking, model versioning
Model Serving	TorchServe / TF Serving + ONNX	Production inference
Vector DB	Pinecone / Qdrant / Weaviate	Embedding storage for semantic search

5.4 Data Engineering & Processing

Component	Technology	Justification
Relational DB	PostgreSQL 15+	ACID compliance, reliability
Document DB	MongoDB 7.x	Flexible schemas, JSON storage
Time-Series DB	TimescaleDB (PostgreSQL extension)	Efficient time-series queries
Cache	Redis 7.x	Session storage, rate limiting, caching
Message Queue	Apache Kafka + RabbitMQ	Event streaming, async processing
Search Engine	Elasticsearch	Full-text search, log aggregation
Object Storage	MinIO / AWS S3	File uploads, model storage
Data Pipeline	Apache Airflow	ETL orchestration, scheduled jobs
Streaming	Apache Flink / Kafka Streams	Real-time data processing

5.5 Blockchain & Security

Component	Technology	Justification
Blockchain	Hyperledger Fabric / Ethereum	Private blockchain for carbon credits
Smart Contracts	Solidity (for Ethereum)	Automated carbon credit verification
Encryption	AES-256 + RSA	Data encryption at rest and in transit
TLS/SSL	Let's Encrypt + Certbot	HTTPS enforcement
Secrets Management	HashiCorp Vault	API keys, credentials
RBAC	Casbin	Policy-based access control
Rate Limiting	Redis + Token Bucket Algorithm	API protection
WAF	ModSecurity	Web application firewall

5.6 DevOps & Infrastructure

Component	Technology	Justification
Containerization	Docker + Docker Compose	Consistent environments
Orchestration	Kubernetes (K8s)	Container orchestration, auto-scaling
CI/CD	GitHub Actions + Jenkins	Automated testing and deployment
Cloud Provider	AWS / GCP / Azure (choose one)	Scalable infrastructure
IaC	Terraform	Infrastructure as code
Monitoring	Prometheus + Grafana	Metrics and dashboards
Logging	ELK Stack (Elasticsearch, Logstash, Kibana)	Centralized logging
Error Tracking	Sentry	Real-time error monitoring
APM	New Relic / Datadog	Application performance monitoring
Load Balancer	Nginx / HAProxy	Traffic distribution
CDN	CloudFlare	Static asset delivery

5.7 Testing & Quality Assurance

Component	Technology	Justification
Unit Testing	Jest + Pytest	JavaScript and Python testing
Integration Testing	Supertest + pytest-asyncio	API testing
E2E Testing	Cypress + Playwright	Browser automation
Load Testing	K6 + Apache JMeter	Performance testing
Security Testing	OWASP ZAP + Snyk	Vulnerability scanning
Code Quality	SonarQube + ESLint + Pylint	Static code analysis
Code Coverage	Codecov + Coverage.py	Coverage tracking

6. DETAILED FEATURE BREAKDOWN

6.1 Core Features (Semester 1 Priority)

A. User Management & Authentication

- Multi-factor authentication (MFA)
- OAuth2.0 social login (Google, GitHub)
- Role-based access control (Individual, Organization Admin, Auditor)
- Profile management with preferences
- Privacy settings and data export (GDPR compliance)

B. Multi-Modal Data Ingestion 1. **Manual Entry** - Interactive forms with smart suggestions - Voice input for quick logging - Bulk upload via CSV/Excel

2. **Receipt/Invoice OCR (Computer Vision)** - Upload photo of bill/receipt - Extract: vendor, items, quantities, prices, dates - Automatic carbon calculation using product database - Support for printed and handwritten receipts - Multi-language support (English, Hindi, regional languages)

3. **Bank Transaction Integration** - Connect bank accounts via Plaid/Yodlee APIs - Automatic categorization using NLP - Transaction-to-carbon mapping - Spend pattern analysis

4. **Calendar Integration** - Google Calendar / Outlook integration - Extract travel data from events - Commute pattern recognition - Meeting-related travel emissions

5. **Email Parsing (NLP)** - Scan emails for flight bookings, e-commerce receipts - Extract structured data - Privacy-preserving local processing

6. IoT Device Integration (Future) - Smart meter data (electricity, gas, water) - Vehicle OBD-II data - Smart home devices

C. Intelligent Carbon Calculation Transport Agent: - Mode detection (car, bus, train, flight, bike, walk) - Distance calculation using Maps API - Vehicle type and fuel efficiency - Carpool factor adjustment - Electric vs ICE comparison - Flight class and distance bands

Energy Agent: - Appliance-level breakdown - Time-of-use analysis - Renewable energy percentage - Grid carbon intensity (location-based) - HVAC optimization recommendations

Food Agent: - Meal image recognition (take photo → identify → carbon score) - Ingredient-level analysis - Local vs imported factor - Organic vs conventional - Meat vs plant-based comparison - Food waste estimation

Procurement Agent: - Product carbon database (100K+ products) - Barcode scanning for instant carbon lookup - Eco-alternatives recommendation - Lifecycle analysis (cradle-to-grave) - Supplier carbon ratings

Waste Agent: - Waste categorization (recyclable, compost, landfill) - Image-based waste classification - Recycling potential score - Circular economy opportunities

D. Advanced Analytics & Predictions Time-Series Forecasting: - 7-day, 30-day, 90-day predictions - Trend analysis (increasing/decreasing) - Seasonal decomposition - Anomaly detection (unusual spikes)

Scenario Analysis: - “What-if” simulations - What if I switch to EV? - What if I go vegetarian? - What if I use public transport? - ROI calculations for green investments - Carbon budget planning

Comparative Analytics: - Personal vs city/country average - Percentile ranking - Historical comparison (year-over-year) - Category-wise breakdown

Predictive Insights: - “You’re on track to exceed your monthly carbon budget by 15%” - “Switching to LED bulbs will save 150kg CO2/year” - “Your diet accounts for 40% of your footprint”

E. Explainable AI Dashboard SHAP Visualizations: - Feature importance for predictions - Waterfall charts showing contribution - Force plots for individual predictions

LIME Explanations: - Local interpretable model approximations - “Why did you get this score?” explanations

Attention Maps (for image/text models): - Highlight which parts of receipt influenced the category - Show which words in transaction description mattered

Decision Trees: - Visual rule-based explanations - “If you reduce car travel by X, carbon drops by Y”

Counterfactual Explanations: - “To reduce your score by 20%, you could...”
- Multiple pathway suggestions

F. Behavioral Intelligence & Nudging Reinforcement Learning-Based Nudges: - Learn user response patterns - Personalized intervention timing - Message tone optimization (encouraging vs factual)

Gamification: - Achievement badges (First 100kg saved, Zero waste week) - Leaderboards (opt-in, privacy-preserving) - Challenges (friends/family competitions) - Streaks and milestones

Social Comparison: - Compare with anonymized similar users - Community averages - Goal-setting based on peer benchmarks

Psychological Profiling: - Segment users (eco-warriors, pragmatists, beginners) - Tailor recommendations - A/B test messaging strategies

G. Recommendation Engine Personalized Suggestions: - Top 3 high-impact actions - Quick wins vs long-term investments - Cost-benefit analysis

Product Alternatives: - Green alternatives for purchased items - Local sourcing options - Second-hand/refurbished recommendations

Lifestyle Optimization: - Diet modifications - Transport mode shifts - Energy efficiency upgrades - Waste reduction strategies

H. Blockchain-Verified Carbon Credits Carbon Credit Tokenization:
- 1 token = 1 kg CO₂ offset - Smart contract-based issuance - Automatic calculation from verified reductions

Offset Marketplace: - Buy verified offsets (reforestation, renewable energy) - Sell your earned credits - Transparent pricing

Audit Trail: - Immutable record of all transactions - Third-party verification hooks - Compliance reporting

I. Reporting & Compliance Automated Reports: - Weekly/monthly/annual summaries - PDF/CSV export - Custom date ranges

Compliance Support: - GHG Protocol aligned calculations - ISO 14064 compatible - Corporate sustainability reporting (CSRD, TCFD)

Audit Mode: - Detailed calculation breakdowns - Data source tracking - Verification checkpoints

6.2 Advanced Features (Semester 2 Priority)

J. Enterprise Module

- Multi-user organizations
- Department-level tracking
- Scope 1, 2, 3 calculations
- Supply chain carbon mapping
- Vendor carbon scorecards
- Carbon pricing integration

K. API Platform

- Public API for third-party integrations
- Webhooks for real-time updates
- Developer portal with docs
- Rate-limited tiers

L. Mobile App

- React Native cross-platform app
- Offline-first architecture
- Camera-based quick logging
- Push notifications

M. Advanced ML Models

- Multi-modal transformers (combining text, image, time-series)
- Federated learning for privacy
- Continual learning for model updates
- Graph neural networks for supply chain

7. DETAILED AI/ML IMPLEMENTATION

7.1 Computer Vision Models

A. Receipt/Invoice OCR Pipeline **Stage 1: Text Detection** - Model: CRAFT (Character Region Awareness for Text Detection) - Input: Image (RGB) - Output: Bounding boxes of text regions

Stage 2: Text Recognition - Model: TrOCR (Transformer-based OCR) - Alternative: EasyOCR + Tesseract ensemble - Input: Cropped text regions - Output: Extracted text

Stage 3: Information Extraction (NER) - Model: BERT-based NER fine-tuned on invoice data - Extract: Vendor, Items, Quantities, Prices, Dates - Post-processing: Regex patterns, rule-based validation

Stage 4: Carbon Mapping - Product database lookup (fuzzy matching) - Category classification if not found - Carbon factor application

Training Dataset: - SROIE (Scanned Receipts OCR and Information Extraction) - Custom annotated Indian receipts (500+ samples) - Data augmentation: rotations, lighting variations, noise

Performance Metrics: - Character accuracy: >95% - Field extraction accuracy: >90% - End-to-end carbon calculation accuracy: >85%

B. Food Image Recognition Model Architecture: - Base: EfficientNet-B4 (pre-trained on ImageNet) - Fine-tuned on food datasets - Output: Top-5 food items with confidence scores

Training Dataset: - Food-101 (101,000 images) - Indian Food Dataset (custom collection, 50+ dishes) - Open Images Dataset (food categories)

Carbon Calculation: - Ingredient database (carbon per 100g) - Portion size estimation using object detection - Cooking method factor

Augmentation: - Color jitter, rotation, zoom - Mixup and CutMix

Inference: - TorchServe deployed model - Latency: <500ms - Accuracy: >85% Top-1, >95% Top-5

C. Waste Classification Model: YOLO v8 - Object detection for waste items - Classes: Plastic, Paper, Metal, Glass, Organic, E-waste, Hazardous - Real-time inference on mobile devices

Dataset: - TACO (Trash Annotations in Context) - Custom waste images (1000+ images)

Output: - Detected items with bounding boxes - Recyclability score - Disposal recommendations

7.2 NLP Models

A. Transaction Categorization Model: FinBERT (Fine-tuned BERT for Finance) - Input: Transaction description - Output: Category (Transport, Food, Energy, Shopping, etc.) - Sub-category classification

Training: - Bank transaction datasets (25,000+ samples) - Manual annotation for carbon-relevant categories - Transfer learning from FinBERT

Emission Calculation: - Category-specific carbon factors - Amount-based scaling - Merchant-specific adjustments

B. Document Analysis Policy Extraction: - Model: LayoutLMv3 (Document understanding) - Extract: Carbon policies, regulations, targets from PDFs - Structured output: JSON

Sentiment Analysis: - Model: RoBERTa fine-tuned on sustainability texts - Analyze: Company sustainability reports - Output: Sentiment scores, commitment detection

Question Answering: - Model: BERT-based QA - User asks: “What’s my biggest carbon source?” - System answers with context and evidence

7.3 Time-Series Forecasting Models

A. Prophet (Meta)

- Used for: Daily/weekly carbon footprint forecasting
- Handles: Seasonality, holidays, missing data
- Uncertainty intervals

B. N-BEATS (Neural Basis Expansion)

- Deep learning for time series
- Pure time-series features (no covariates)
- Interpretable and accurate

C. Temporal Fusion Transformer (TFT)

- Attention-based architecture
- Multi-horizon forecasting
- Handles multiple time series simultaneously
- Incorporates known future covariates (weather, events)

Ensemble Approach: - Combine predictions from all three - Weighted average based on recent performance - Fallback to simpler models if data is sparse

7.4 Reinforcement Learning for Behavioral Agent

A. Nudge Generation Environment: - State: User profile, historical behavior, current carbon level - Action: Send nudge (type, timing, message) - Reward: User engagement + carbon reduction

Algorithm: Deep Q-Network (DQN) - Neural network approximates Q-values - -greedy exploration - Experience replay

Personalization: - User segments learned through clustering - A/B testing for message effectiveness - Multi-armed bandit for optimal timing

B. Route Optimization Algorithm: Proximal Policy Optimization (PPO) - State: Current location, destination, time, weather - Action: Suggest transport mode - Reward: (Carbon saved) - (Time penalty) - (Cost penalty)

Integration: - Google Maps API for routes - Real-time carbon intensity of grid (for EV charging) - User preferences (comfort, speed, cost)

7.5 Explainable AI Implementation

A. SHAP (SHapley Additive exPlanations)

```
import shap

# For tree-based models
explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X)

# For deep learning models
explainer = shap.DeepExplainer(model, background_data)
shap_values = explainer.shap_values(X)

# Visualization
shap.summary_plot(shap_values, X)
shap.force_plot(base_value, shap_values[0], X.iloc[0])
```

B. LIME

```
from lime.lime_tabular import LimeTabularExplainer

explainer = LimeTabularExplainer(
    training_data=X_train,
    feature_names=feature_names,
    class_names=['Low', 'Medium', 'High'],
    mode='classification'
)

explanation = explainer.explain_instance(
    data_row=X_test[0],
    predict_fn=model.predict_proba
)
```

C. Attention Visualization (for Transformers)

```
import torch
from transformers import BertModel, BertTokenizer

# Extract attention weights
outputs = model(input_ids, attention_mask, output_attentions=True)
attention = outputs.attentions # Tuple of attention matrices

# Visualize using BertViz
from bertviz import head_view
head_view(attention, tokens)
```

8. DATABASE SCHEMA DESIGN

8.1 PostgreSQL (Relational Data)

Users Table

```
CREATE TABLE users (  
    user_id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
    email VARCHAR(255) UNIQUE NOT NULL,  
    password_hash VARCHAR(255) NOT NULL,  
    full_name VARCHAR(255),  
    role VARCHAR(50) DEFAULT 'individual', -- individual, org_admin, auditor  
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
    is_verified BOOLEAN DEFAULT FALSE,  
    preferences JSONB,  
    location GEOGRAPHY(POINT),  
    timezone VARCHAR(50)  
);  
  
CREATE INDEX idx_users_email ON users(email);  
CREATE INDEX idx_users_location ON users USING GIST(location);
```

Carbon_Transactions Table

```
CREATE TABLE carbon_transactions (  
    transaction_id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
    user_id UUID REFERENCES users(user_id),  
    category VARCHAR(50) NOT NULL, -- transport, energy, food, waste, procurement  
    sub_category VARCHAR(100),  
    amount DECIMAL(10,2),  
    carbon_kg DECIMAL(10,3) NOT NULL,  
    confidence_score DECIMAL(3,2), -- 0.00 to 1.00  
    data_source VARCHAR(50), -- manual, ocr, bank, api  
    source_metadata JSONB,  
    transaction_date DATE NOT NULL,  
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
    is_verified BOOLEAN DEFAULT FALSE,  
    verification_method VARCHAR(50),  
    blockchain_tx_hash VARCHAR(255) -- Link to blockchain  
);  
  
CREATE INDEX idx_carbon_user_date ON carbon_transactions(user_id, transaction_date);  
CREATE INDEX idx_carbon_category ON carbon_transactions(category);
```

Carbon_Credits Table

```

CREATE TABLE carbon_credits (
    credit_id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    user_id UUID REFERENCES users(user_id),
    amount_kg DECIMAL(10,3) NOT NULL,
    credit_type VARCHAR(50), -- earned, purchased, offset
    source VARCHAR(100), -- reforestation, renewable_energy, etc.
    price_per_kg DECIMAL(10,2),
    blockchain_token_id VARCHAR(255),
    issued_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    expires_at TIMESTAMP,
    status VARCHAR(50) DEFAULT 'active' -- active, retired, transferred
);

```

AI_Models Table

```

CREATE TABLE ai_models (
    model_id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    model_name VARCHAR(255) NOT NULL,
    model_type VARCHAR(50), -- cv, nlp, timeseries, rl
    version VARCHAR(50),
    architecture VARCHAR(100),
    trained_at TIMESTAMP,
    metrics JSONB, -- accuracy, precision, recall, etc.
    model_path VARCHAR(500),
    is_active BOOLEAN DEFAULT TRUE,
    created_by UUID REFERENCES users(user_id)
);

```

8.2 MongoDB (Document Store)

receipts Collection

```

{
  _id: ObjectId(),
  userId: UUID,
  uploadedAt: ISODate(),
  imageUrl: String,
  ocrText: String,
  extractedData: {
    vendor: String,
    items: [
      {
        name: String,
        quantity: Number,
        price: Number,
        carbonKg: Number
      }
    ]
  }
}

```

```

    ],
    totalAmount: Number,
    totalCarbon: Number,
    date: ISODate()
  },
  processingStatus: String, // pending, success, failed
  confidence: Number
}

```

nudges Collection

```

{
  _id: ObjectId(),
  userId: UUID,
  nudgeType: String, // behavioral, informational, gamification
  message: String,
  sentAt: ISODate(),
  channel: String, // push, email, in-app
  userAction: String, // clicked, dismissed, ignored
  actionTimestamp: ISODate(),
  rlState: Object, // RL agent state
  rlAction: String,
  reward: Number
}

```

predictions Collection

```

{
  _id: ObjectId(),
  userId: UUID,
  predictionType: String, // daily, weekly, monthly
  forecastedValues: [
    {
      date: ISODate(),
      carbonKg: Number,
      lowerBound: Number,
      upperBound: Number
    }
  ],
  modelUsed: String,
  generatedAt: ISODate(),
  accuracy: Number // Compared to actuals later
}

```

8.3 TimescaleDB (Time-Series Data)

```
CREATE TABLE carbon_timeseries (  
    time TIMESTAMPTZ NOT NULL,  
    user_id UUID NOT NULL,  
    category VARCHAR(50),  
    carbon_kg DECIMAL(10,3),  
    metadata JSONB  
);  
  
SELECT create_hypertable('carbon_timeseries', 'time');  
  
CREATE INDEX idx_timeseries_user ON carbon_timeseries (user_id, time DESC);
```

8.4 Redis (Caching & Sessions)

Session Storage:

Key: session:{sessionId}
Value: {userId, role, expiresAt}
TTL: 7 days

Rate Limiting:

Key: rate_limit:{userId}:{endpoint}
Value: request_count
TTL: 1 hour

Model Predictions Cache:

Key: prediction:{userId}:{date}
Value: {carbonKg, confidence}
TTL: 24 hours

8.5 Vector Database (Pinecone/Qdrant)

Embeddings for Semantic Search: - Receipt text embeddings (find similar purchases) - Product embeddings (recommend alternatives) - User behavior embeddings (find similar users)

Example: Store receipt embedding

```
index.upsert(vectors=[  
    {  
        'id': receipt_id,  
        'values': embedding_vector, # 768-dim from BERT  
        'metadata': {  
            'userId': user_id,  
            'category': 'food',  
            'carbonKg': 5.2  
        }  
    }  
])
```

```

    }
  })

  # Query similar receipts
  results = index.query(
    vector=query_embedding,
    top_k=10,
    filter={'category': 'food'}
  )

```

9. API DESIGN

9.1 RESTful API Endpoints

Authentication

```

POST /api/v1/auth/register
POST /api/v1/auth/login
POST /api/v1/auth/logout
POST /api/v1/auth/refresh-token
POST /api/v1/auth/forgot-password
POST /api/v1/auth/reset-password
GET /api/v1/auth/verify-email/:token

```

User Management

```

GET /api/v1/users/me
PUT /api/v1/users/me
DELETE /api/v1/users/me
GET /api/v1/users/me/preferences
PUT /api/v1/users/me/preferences

```

Carbon Tracking

```

GET /api/v1/carbon/transactions
POST /api/v1/carbon/transactions
GET /api/v1/carbon/transactions/:id
PUT /api/v1/carbon/transactions/:id
DELETE /api/v1/carbon/transactions/:id
GET /api/v1/carbon/summary?period=weekly
GET /api/v1/carbon/category-breakdown
GET /api/v1/carbon/trends

```

OCR & Computer Vision

```

POST /api/v1/cv/upload-receipt

```

```
GET    /api/v1/cv/receipt/:id
POST   /api/v1/cv/food-recognition
POST   /api/v1/cv/waste-classification
```

Predictions

```
GET    /api/v1/predictions/forecast?horizon=30
GET    /api/v1/predictions/anomalies
POST   /api/v1/predictions/scenario
```

Recommendations

```
GET    /api/v1/recommendations/actions
GET    /api/v1/recommendations/products/:productId/alternatives
GET    /api/v1/recommendations/personalized
```

Carbon Credits

```
GET    /api/v1/credits/balance
GET    /api/v1/credits/transactions
POST   /api/v1/credits/offset
POST   /api/v1/credits/buy
POST   /api/v1/credits/sell
GET    /api/v1/credits/marketplace
```

Analytics

```
GET    /api/v1/analytics/compare?type=city
GET    /api/v1/analytics/leaderboard
GET    /api/v1/analytics/achievements
GET    /api/v1/analytics/impact
```

Reports

```
GET    /api/v1/reports/generate?period=monthly&format=pdf
GET    /api/v1/reports/compliance
GET    /api/v1/reports/audit-trail
```

9.2 GraphQL Schema (Alternative for Complex Queries)

```
type User {
  id: ID!
  email: String!
  fullName: String
  role: UserRole!
  carbonFootprint: CarbonFootprint!
  predictions: [Prediction!]!
  recommendations: [Recommendation!]!
```

```

    credits: CarbonCredits!
  }

  type CarbonFootprint {
    totalKg: Float!
    byCategory: [CategoryBreakdown!]!
    trend: TrendData!
    percentile: Int
  }

  type CategoryBreakdown {
    category: Category!
    carbonKg: Float!
    percentage: Float!
    transactions: [Transaction!]!
  }

  type Prediction {
    date: Date!
    carbonKg: Float!
    confidence: Float!
    explanation: String
  }

  type Query {
    me: User!
    carbonSummary(period: Period!): CarbonFootprint!
    forecast(horizon: Int!): [Prediction!]!
    recommendations(limit: Int): [Recommendation!]!
  }

  type Mutation {
    createTransaction(input: TransactionInput!): Transaction!
    uploadReceipt(file: Upload!): Receipt!
    offsetCarbon(amount: Float!): CarbonOffset!
  }

```

10. TWO-SEMESTER IMPLEMENTATION PLAN

SEMESTER 1 (Months 1-6): FOUNDATION & CORE FEATURES

Month 1: Project Setup & Architecture **Week 1-2: Infrastructure Setup** - Set up version control (GitHub/GitLab) - Initialize monorepo structure (Nx/Turborepo) - Set up development environments - Configure Docker containers - Set up CI/CD pipelines (GitHub Actions) - Initialize databases

(PostgreSQL, MongoDB, Redis) - Set up monitoring (Prometheus + Grafana)

Tasks Division: - Member 1: Frontend setup + Docker - Member 2: Backend setup + Databases - Member 3: CI/CD + Monitoring

Week 3-4: Core Backend Development - User authentication system (Auth0 + JWT) - Basic CRUD APIs for users - Database migrations - API documentation (Swagger) - Unit tests for auth

Deliverable: Working authentication system with API docs

Month 2: Data Ingestion & OCR Week 1-2: Manual Data Entry & Basic Calculation - Frontend forms for carbon tracking - Backend calculation engine - Carbon factor database setup - Category-wise calculations - Basic dashboard

Week 3-4: Receipt OCR Implementation - Tesseract + EasyOCR integration - Image preprocessing pipeline - Text extraction and parsing - NER for information extraction - Integration with carbon calculation

Tasks Division: - Member 1: Frontend forms + Dashboard - Member 2: Calculation engine + Carbon DB - Member 3: OCR pipeline + NER

Deliverable: Working manual entry + OCR-based carbon tracking

Month 3: Computer Vision & NLP Week 1-2: Food Image Recognition - EfficientNet model training - Food-101 dataset fine-tuning - Indian food dataset collection - Model deployment (TorchServe) - Frontend integration

Week 3-4: Transaction Categorization (NLP) - FinBERT fine-tuning - Transaction dataset preparation - Model training and evaluation - API endpoint creation - Testing with real transactions

Tasks Division: - Member 1: Food recognition frontend - Member 2: Food model training + deployment - Member 3: NLP model + API

Deliverable: Food recognition + transaction categorization working

Month 4: Time-Series Forecasting & Analytics Week 1-2: Forecasting Models - Prophet implementation - N-BEATS implementation - Temporal Fusion Transformer (TFT) - Ensemble approach - Backtesting and evaluation

Week 3-4: Analytics Dashboard - Time-series visualization (D3.js) - Category breakdown charts - Trend analysis - Comparative analytics (vs city/country average) - Anomaly detection alerts

Tasks Division: - Member 1: Forecasting models - Member 2: Analytics backend APIs - Member 3: Visualization dashboard

Deliverable: Predictive analytics dashboard

Month 5: Explainable AI & Recommendations Week 1-2: XAI Implementation - SHAP integration for all models - LIME for local explanations - Attention visualization for NLP/CV - Frontend XAI dashboard - Interactive explanation widgets

Week 3-4: Recommendation Engine - Collaborative filtering for users - Content-based product recommendations - Hybrid recommendation system - Top-N action suggestions - A/B testing framework

Tasks Division: - Member 1: XAI frontend - Member 2: SHAP/LIME backend - Member 3: Recommendation engine

Deliverable: Explainable predictions + personalized recommendations

Month 6: Testing, Security & Semester 1 Deployment Week 1-2: Comprehensive Testing - Unit test coverage >80% - Integration tests for all APIs - E2E tests (Cypress) - Load testing (K6) - Security testing (OWASP ZAP)

Week 3: Security Hardening - Penetration testing - Vulnerability fixes - Rate limiting implementation - Data encryption verification - GDPR compliance audit

Week 4: Deployment & Documentation - Production deployment (AWS/GCP) - User documentation - API documentation - Video demo creation - Project report (Semester 1)

Tasks Division: - Member 1: Testing + fixes - Member 2: Security + deployment - Member 3: Documentation + demo

Deliverable: Deployed platform + comprehensive documentation

SEMESTER 2 (Months 7-12): ADVANCED FEATURES & SCALING

Month 7: Behavioral AI & Gamification Week 1-2: Reinforcement Learning Agent - DQN implementation for nudging - User behavior dataset creation - Reward function design - Training and evaluation - Integration with notification system

Week 3-4: Gamification System - Achievement system - Badge design - Leaderboard (privacy-preserving) - Challenge creation - Streak tracking

Deliverable: Intelligent nudging + gamification

Month 8: Blockchain Integration **Week 1-2: Smart Contracts** - Hyperledger Fabric setup - Carbon credit smart contracts - Token issuance logic - Verification mechanisms - Testing on testnet

Week 3-4: Marketplace - Frontend marketplace UI - Buy/sell functionality - Pricing engine - Transaction history - Blockchain explorer integration

Deliverable: Blockchain-verified carbon credits + marketplace

Month 9: Enterprise Features **Week 1-2: Multi-User Organizations** - Organization schema - Department-level tracking - Scope 1, 2, 3 calculations - Admin panel - User role management

Week 3-4: Supply Chain Mapping - Vendor carbon scorecards - Supply chain visualization (Graph) - Upstream/downstream emissions - Integration with procurement systems

Deliverable: Enterprise module

Month 10: Mobile App & Advanced ML **Week 1-2: React Native App** - Cross-platform mobile app - Offline-first architecture - Camera integration - Push notifications - App store deployment

Week 3-4: Advanced ML Models - Multi-modal transformer (text + image + time-series) - Transfer learning across domains - Federated learning setup (privacy) - Continual learning pipeline

Deliverable: Mobile app + advanced ML

Month 11: Scaling & Optimization **Week 1-2: Performance Optimization** - Database query optimization - Caching strategy refinement - API response time <200ms - Image optimization (WebP) - Code splitting and lazy loading

Week 3-4: Kubernetes & Auto-Scaling - K8s cluster setup - Horizontal pod autoscaling - Load balancer configuration - Multi-region deployment - Disaster recovery

Deliverable: Highly optimized, scalable platform

Month 12: Final Testing & Launch **Week 1: Comprehensive Testing**
- Regression testing - User acceptance testing (UAT) - Beta user feedback - Bug fixes

Week 2: Documentation & Compliance - User manual - Admin guide - API reference - Security audit report - Privacy policy & ToS

Week 3: Marketing & Deployment - Landing page - Video tutorials - Social media launch - Press release - Production deployment

Week 4: Project Presentation - Final project report - Live demonstration - Presentation to faculty - Q&A preparation

Deliverable: Production-ready platform + comprehensive documentation

11. EVALUATION METRICS & KPIs

11.1 Technical Metrics

Model Performance

- **OCR Accuracy:** >95% character-level accuracy
- **Food Recognition:** >85% Top-1, >95% Top-5
- **Transaction Categorization:** >90% F1-score
- **Forecasting:** MAPE <10% for 7-day, <15% for 30-day
- **Waste Classification:** >88% mAP

System Performance

- **API Response Time:** <200ms for 95th percentile
- **Dashboard Load Time:** <2 seconds
- **Concurrent Users:** Support 10,000+ simultaneously
- **Uptime:** 99.9% SLA
- **Data Processing:** 1000+ receipts/minute

Security & Compliance

- **Encryption:** AES-256 for data at rest, TLS 1.3 for transit
- **Vulnerability Score:** 0 critical, 0 high (from scans)
- **GDPR Compliance:** 100% data export/deletion capability
- **Authentication:** MFA adoption >60%

11.2 User Metrics

- **User Engagement:** DAU/MAU ratio >40%
- **Retention:** 60% 30-day retention

- **Carbon Reduction:** Average 15% reduction in first 90 days
- **Data Quality:** >80% verified transactions
- **User Satisfaction:** NPS >50

11.3 Business Metrics

- **Carbon Credits Traded:** \$10,000+ in first year
- **Enterprise Clients:** 5+ pilot organizations
- **API Integrations:** 10+ third-party apps
- **Cost per User:** <\$2/month at scale

12. RISK MITIGATION

12.1 Technical Risks

Risk	Probability	Impact	Mitigation
Model accuracy below target	Medium	High	Ensemble methods, continuous retraining, human-in-loop validation
Scalability issues	Medium	High	K8s auto-scaling, load testing from Month 1, CDN for static assets
Data privacy breach	Low	Critical	End-to-end encryption, regular security audits, bug bounty program
Third-party API failures	Medium	Medium	Fallback mechanisms, caching, multiple providers
Blockchain transaction delays	Low	Medium	Layer-2 solutions, off-chain aggregation

12.2 Project Risks

Risk	Probability	Impact	Mitigation
Team member unavailability	Medium	High	Cross-training, documentation, pair programming
Scope creep	High	Medium	Strict sprint planning, MVP first, feature freezes
Technology learning curve	Medium	Medium	Proof-of-concepts first, online courses, mentorship
Timeline delays	Medium	High	Buffer time (10%), weekly progress reviews, agile methodology

13. INNOVATION & RESEARCH CONTRIBUTIONS

13.1 Novel Approaches

1. **Multi-Agentic Architecture for Carbon Tracking**
 - First platform to use specialized AI agents
 - Agent coordination using message passing
 - Domain-specific expertise
2. **Cross-Modal Learning**
 - Combining text (receipts) + images (products) + time-series (patterns)
 - Unified embedding space
 - Transfer learning across modalities
3. **Behavioral RL for Sustainability**
 - Novel reward function balancing engagement and impact
 - Personalized intervention strategies
 - Ethical AI considerations
4. **Blockchain + AI Integration**
 - Smart contracts for carbon credit verification
 - AI-driven pricing mechanisms
 - Transparent audit trails
5. **Federated Learning for Privacy**
 - Train models without centralizing data
 - Enterprise collaboration without data sharing
 - Differential privacy guarantees

13.2 Potential Publications

1. “Multi-Agentic AI for Personalized Carbon Footprint Monitoring”
 - Target: IEEE Transactions on Sustainable Computing
 2. “Explainable Deep Learning for Carbon Emission Prediction”
 - Target: ACM Conference on Fairness, Accountability, and Transparency (ACM FAccT)
 3. “Behavioral Reinforcement Learning for Environmental Nudging”
 - Target: AAAI Conference on Artificial Intelligence
 4. “Blockchain-Verified Carbon Credits: A Decentralized Approach”
 - Target: IEEE International Conference on Blockchain
-

14. COMMERCIALIZATION POTENTIAL

14.1 Business Model

Freemium Model: - **Free Tier:** Basic tracking, manual entry, limited predictions - **Pro Tier (\$9.99/month):** OCR, advanced analytics, carbon credits, priority support - **Enterprise Tier (Custom pricing):** Multi-user, API access, white-label, SLA

Revenue Streams: - Subscription fees - Carbon credit marketplace (5% transaction fee) - API usage (pay-per-call) - Data insights (anonymized, aggregated) - Consulting services

14.2 Go-to-Market Strategy

Phase 1 (Semester 2): Beta launch with 100 users **Phase 2 (Post-graduation):** Public launch, target 10,000 users in 6 months **Phase 3:** Enterprise sales, partnerships with sustainability consultants **Phase 4:** International expansion

14.3 Funding Opportunities

- Startup India Seed Fund
 - University incubation programs
 - Climate tech accelerators (Y Combinator, Techstars)
 - Government grants (UNDP, GIZ)
 - Angel investors in sustainability space
-

15. SUSTAINABILITY & ETHICS

15.1 Platform Sustainability

- **Green Hosting:** Use carbon-neutral cloud providers (Google Cloud with 100% renewable energy match)
- **Efficient Computing:** Optimize models for inference (quantization, pruning)
- **Offset Own Footprint:** Track and offset platform's carbon emissions

15.2 Ethical Considerations

- **Privacy First:** User data never sold, transparent data usage
 - **Bias Mitigation:** Test models for demographic bias, diverse training data
 - **Accessibility:** WCAG 2.1 AA compliance, multilingual support
 - **Inclusivity:** Consider low-income users, avoid shaming language
 - **Transparency:** Open-source non-proprietary components
-

16. COMPETITIVE ANALYSIS

16.1 Existing Solutions

Competitor	Strengths	Weaknesses	Our Advantage
Carbon Footprint App	Simple, mobile-first	Manual entry only, no AI	Multi-modal AI, automation
Joro	Bank integration	US-only, basic analytics	Global, advanced ML, XAI
Pawprint	Gamification	Limited data sources	Multi-agentic, enterprise features
Capture	Enterprise focus	No individual users	Both B2C and B2B
Patch	Carbon offsets	No tracking	Integrated tracking + offsets

16.2 Differentiation

1. **Only platform with multi-agentic AI** specializing in different emission sources
2. **Strongest XAI capabilities** for trust and compliance
3. **Blockchain verification** for carbon credits (unique)
4. **Advanced computer vision** (receipt OCR, food recognition, waste classification)
5. **Behavioral RL** for personalized nudging

6. Enterprise + Individual dual focus

17. SUCCESS CRITERIA

17.1 Semester 1 Success

Fully functional web platform deployed OCR accuracy >90% Time-series forecasting working XAI dashboard operational 50+ beta users testing Comprehensive documentation >80% code coverage Security audit passed

17.2 Semester 2 Success

Mobile app launched Blockchain integration complete Enterprise module functional 500+ active users Average 15% carbon reduction demonstrated API published and documented 1+ conference paper submitted Commercialization plan ready

17.3 Post-Project Success

10,000+ users in 6 months 5+ enterprise clients Funding secured Profitable unit economics Measurable real-world carbon reductions (100+ tons CO2 offset)

18. TEAM RESPONSIBILITIES

18.1 Recommended Division

Member 1: Frontend & UI/UX - React application development - Dashboard design - Data visualization (D3.js) - Mobile app (React Native) - User testing

Member 2: Backend & Infrastructure - Node.js microservices - Database design and optimization - API development - DevOps and deployment - Security implementation

Member 3: AI/ML & Data Science - Model training and optimization - Computer vision pipeline - NLP models - Time-series forecasting - XAI implementation

Shared Responsibilities: - Architecture decisions (weekly meetings) - Code reviews - Testing - Documentation - Presentations

18.2 Collaboration Tools

- **Version Control:** GitHub with branch protection
- **Project Management:** Jira / Linear

- **Communication:** Slack / Discord
 - **Documentation:** Notion / Confluence
 - **Design:** Figma
 - **API Testing:** Postman
 - **Monitoring:** Sentry + Datadog
-

19. LEARNING RESOURCES

19.1 Required Skills & Learning Paths

For All Members: - Git & GitHub Advanced - Docker & Kubernetes basics - Microservices architecture - RESTful API design - Testing best practices

Frontend Developer: - React Deep Dive (Udemy: Maximilian Schwarzmüller) - TypeScript Masterclass - D3.js for Data Visualization - React Native (if doing mobile)

Backend Developer: - Node.js Advanced Patterns - PostgreSQL Performance Tuning - MongoDB University - Kafka & message queues - OAuth2.0 & JWT

AI/ML Developer: - Fast.ai Practical Deep Learning - Hugging Face Transformers Course - PyTorch Lightning - MLOps (DeepLearning.AI) - SHAP & LIME tutorials

19.2 Recommended Certifications

- AWS Certified Solutions Architect (for deployment)
 - TensorFlow Developer Certificate
 - Certified Kubernetes Administrator (CKA)
-

20. FINAL THOUGHTS & VISION

CarbonSense AI is not just a college project—it's a platform with real-world impact potential. By combining cutting-edge AI/ML techniques with a genuine sustainability mission, this project can:

1. **Demonstrate Technical Excellence:** Showcasing mastery of ML, DL, NLP, CV, and software engineering
2. **Create Real Impact:** Help thousands reduce their carbon footprint
3. **Enable Research:** Publish papers, contribute to climate tech knowledge
4. **Launch a Startup:** Commercialize post-graduation
5. **Build a Portfolio:** Stand out in job applications and graduate school admissions

Key to Success: - Start simple, iterate fast (Agile) - Test early and often - Focus on user experience - Document everything - Don't over-engineer, but plan

for scale - Have fun building something meaningful!

Remember: The best projects are those that solve real problems. Climate change is the defining challenge of our generation. This platform can be your contribution to the solution.

Good luck! Build something amazing!